

23012064_dsm2

November 5, 2024

Name: Abdallah Saber

ID: 23012064

```
[ ]: # Importing Libraries
import numpy as np
import pandas as pd

import time, warnings
import datetime as dt

#visualizations
import matplotlib.pyplot as plt
from pandas.plotting import scatter_matrix
import seaborn as sns
```

1 Get the Data

```
[2]: #load the dataset
retail_df = pd.read_csv('./data.csv',encoding="ISO-8859-1",dtype={'CustomerID':_
↳str,'InvoiceID': str})
retail_df.head()
```

```
[2]: InvoiceNo StockCode Description Quantity \
0 536365 85123A WHITE HANGING HEART T-LIGHT HOLDER 6
1 536365 71053 WHITE METAL LANTERN 6
2 536365 84406B CREAM CUPID HEARTS COAT HANGER 8
3 536365 84029G KNITTED UNION FLAG HOT WATER BOTTLE 6
4 536365 84029E RED WOOLLY HOTTIE WHITE HEART. 6
```

```
InvoiceDate UnitPrice CustomerID Country
0 12/1/2010 8:26 2.55 17850 United Kingdom
1 12/1/2010 8:26 3.39 17850 United Kingdom
2 12/1/2010 8:26 2.75 17850 United Kingdom
3 12/1/2010 8:26 3.39 17850 United Kingdom
4 12/1/2010 8:26 3.39 17850 United Kingdom
```

2 Prepare the Data

restrict the data to only United Kingdom customers

```
[3]: retail_uk = retail_df[retail_df['Country']=='United Kingdom']  
     #check the shape  
     retail_uk.shape
```

```
[3]: (495478, 8)
```

```
[4]: #remove canceled orders  
     retail_uk = retail_uk[retail_uk['Quantity']>0]  
     retail_uk.shape
```

```
[4]: (486286, 8)
```

```
[5]: #remove rows where customerID are NA  
     retail_uk.dropna(subset=['CustomerID'],how='all',inplace=True)  
     retail_uk.shape
```

```
[5]: (354345, 8)
```

```
[6]: #restrict the data to one full year because it's better to use a metric per_  
     ↪Months or Years in RFM  
     retail_uk = retail_uk[retail_uk['InvoiceDate']>= "2010-12-09"]  
     retail_uk.shape
```

```
[6]: (176137, 8)
```

```
[7]: print("Summary..  
     #exploring the unique values of each attribute  
     print("Number of transactions: ", retail_uk['InvoiceNo'].nunique()  
     print("Number of products bought: ",retail_uk['StockCode'].nunique()  
     print("Number of customers:", retail_uk['CustomerID'].nunique() )  
     print("Percentage of customers NA: ", round(retail_uk['CustomerID'].isnull().  
     ↪sum() * 100 / len(retail_df),2),"%") )
```

Summary..

Number of transactions: 8789

Number of products bought: 3294

Number of customers: 2864

Percentage of customers NA: 0.0 %

3 RFM Analysis

3.1 Recency

```
[8]: #last date available in our dataset
retail_uk['InvoiceDate'].max()
```

```
[8]: '9/9/2011 9:52'
```

The last date we have is 2011-12-09 so we will use it as reference.

```
[9]: now = dt.date(2011,12,9)
print(now)
```

2011-12-09

```
[10]: #create a new column called date which contains the date of invoice only
retail_uk['date'] = pd.DatetimeIndex(retail_uk['InvoiceDate']).date
```

```
[11]: retail_uk.head()
```

```
[11]:
```

	InvoiceNo	StockCode	Description	Quantity	\
105335	545220	21955	DOORMAT UNION JACK GUNS AND ROSES	2	
105336	545220	48194	DOORMAT HEARTS	2	
105337	545220	22556	PLASTERS IN TIN CIRCUS PARADE	12	
105338	545220	22139	RETROSPOT TEA SET CERAMIC 11 PC	3	
105339	545220	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	4	

	InvoiceDate	UnitPrice	CustomerID	Country	date
105335	3/1/2011 8:30	7.95	14620	United Kingdom	2011-03-01
105336	3/1/2011 8:30	7.95	14620	United Kingdom	2011-03-01
105337	3/1/2011 8:30	1.65	14620	United Kingdom	2011-03-01
105338	3/1/2011 8:30	4.95	14620	United Kingdom	2011-03-01
105339	3/1/2011 8:30	3.75	14620	United Kingdom	2011-03-01

```
[12]: #group by customers and check last date of purshace
recency_df = retail_uk.groupby(by='CustomerID', as_index=False)['date'].max()
recency_df.columns = ['CustomerID', 'LastPurshaceDate']
recency_df.head()
```

```
[12]:
```

	CustomerID	LastPurshaceDate
0	12747	2011-08-22
1	12748	2011-09-30
2	12749	2011-08-01
3	12820	2011-09-26
4	12821	2011-05-09

```
[13]: #calculate recency
```

```
recency_df['Recency'] = recency_df['LastPurshaceDate'].apply(lambda x: (now - x).days)
```

```
[14]: recency_df.head()
```

```
[14]:  CustomerID  LastPurshaceDate  Recency
0      12747      2011-08-22      109
1      12748      2011-09-30       70
2      12749      2011-08-01      130
3      12820      2011-09-26       74
4      12821      2011-05-09      214
```

```
[15]: #drop LastPurchaseDate as we don't need it anymore
recency_df.drop('LastPurshaceDate',axis=1,inplace=True)
```

3.2 Frequency

```
[16]: # drop duplicates
retail_uk_copy = retail_uk
retail_uk_copy.drop_duplicates(subset=['InvoiceNo', 'CustomerID'],
                               keep="first", inplace=True)
#calculate frequency of purchases
frequency_df = retail_uk_copy.groupby(by=['CustomerID'],
                                       as_index=False)['InvoiceNo'].count()
frequency_df.columns = ['CustomerID', 'Frequency']
frequency_df.head()
```

```
[16]:  CustomerID  Frequency
0      12747          5
1      12748         96
2      12749          3
3      12820          1
4      12821          1
```

3.3 Monetary

```
[17]: #create column total cost
retail_uk['TotalCost'] = retail_uk['Quantity'] * retail_uk['UnitPrice']
```

```
[18]: monetary_df = retail_uk.groupby(by='CustomerID',as_index=False).
      agg({'TotalCost': 'sum'})
monetary_df.columns = ['CustomerID', 'Monetary']
monetary_df.head()
```

```
[18]:  CustomerID  Monetary
0      12747      191.85
1      12748     1054.43
```

2	12749	67.00
3	12820	15.00
4	12821	19.92

3.4 Create RFM Table

```
[19]: #merge recency dataframe with frequency dataframe
temp_df = recency_df.merge(frequency_df,on='CustomerID')
temp_df.head()
```

```
[19]:  CustomerID  Recency  Frequency
0      12747      109         5
1      12748       70        96
2      12749      130         3
3      12820       74         1
4      12821      214         1
```

```
[20]: #merge with monetary dataframe to get a table with the 3 columns
rfm_df = temp_df.merge(monetary_df,on='CustomerID')
#use CustomerID as index
rfm_df.set_index('CustomerID',inplace=True)
#check the head
rfm_df.head()
```

```
[20]:           Recency  Frequency  Monetary
CustomerID
12747           109         5    191.85
12748           70        96   1054.43
12749          130         3     67.00
12820           74         1     15.00
12821          214         1     19.92
```

3.5 RFM Table Correctness verification

```
[21]: retail_uk[retail_uk['CustomerID']=='12820']
```

```
[21]:      InvoiceNo  StockCode      Description  Quantity \
360567    568236    23328  SET 6 SCHOOL MILK BOTTLES IN CRATE      4

      InvoiceDate  UnitPrice  CustomerID      Country      date \
360567  9/26/2011 11:49      3.75      12820  United Kingdom  2011-09-26

      TotalCost
360567      15.0
```

```
[22]: (now - dt.date(2011,9,26)).days == 74
```

```
[22]: True
```

3.5.1 RFM Quartiles

```
[23]: quantiles = rfm_df.quantile(q=[0.25,0.5,0.75])
      quantiles
```

```
[23]:
```

	Recency	Frequency	Monetary
0.25	85.0	1.0	16.35
0.50	119.0	2.0	35.40
0.75	183.0	3.0	92.42

```
[24]: quantiles.to_dict()
```

```
[24]: {'Recency': {0.25: 85.0, 0.5: 119.0, 0.75: 183.0},
      'Frequency': {0.25: 1.0, 0.5: 2.0, 0.75: 3.0},
      'Monetary': {0.25: 16.35, 0.5: 35.400000000000006, 0.75: 92.42}}
```

3.5.2 Creation of RFM Segments

We will create two segmentation classes since, high recency is bad, while high frequency and monetary value is good.

```
[ ]: # Arguments (x = value, p = recency, monetary_value, frequency, d = quartiles_
      ↪dict)
def RScore(x,p,d):
    if x <= d[p][0.25]:
        return 4
    elif x <= d[p][0.50]:
        return 3
    elif x <= d[p][0.75]:
        return 2
    else:
        return 1

# Arguments (x = value, p = recency, monetary_value, frequency, k = quartiles_
      ↪dict)
def FMScore(x,p,d):
    if x <= d[p][0.25]:
        return 1
    elif x <= d[p][0.50]:
        return 2
    elif x <= d[p][0.75]:
        return 3
    else:
        return 4
```

```
[26]: #create rfm segmentation table
      rfm_segmentation = rfm_df
```

```
rfm_segmentation['R_Quartile'] = rfm_segmentation['Recency'].apply(RScore,
↪args=('Recency',quantiles,))
rfm_segmentation['F_Quartile'] = rfm_segmentation['Frequency'].apply(FMScore,
↪args=('Frequency',quantiles,))
rfm_segmentation['M_Quartile'] = rfm_segmentation['Monetary'].apply(FMScore,
↪args=('Monetary',quantiles,))
```

```
[27]: rfm_segmentation.head()
```

```
[27]:
```

	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile
CustomerID						
12747	109	5	191.85	3	4	4
12748	70	96	1054.43	4	4	4
12749	130	3	67.00	2	3	3
12820	74	1	15.00	4	1	1
12821	214	1	19.92	1	1	2

Get the RFM Segment value

```
[28]: rfm_segmentation['RFMScore'] = rfm_segmentation.R_Quartile.map(str) \
      + rfm_segmentation.F_Quartile.map(str) \
      + rfm_segmentation.M_Quartile.map(str)
rfm_segmentation.head()
```

```
[28]:
```

	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	\
CustomerID							
12747	109	5	191.85	3	4	4	
12748	70	96	1054.43	4	4	4	
12749	130	3	67.00	2	3	3	
12820	74	1	15.00	4	1	1	
12821	214	1	19.92	1	1	2	

```
RFMScore
CustomerID
12747      344
12748      444
12749      233
12820      411
12821      112
```

(best customers)

```
[29]: rfm_segmentation[rfm_segmentation['RFMScore']=='444'].sort_values('Monetary',
↪ascending=False).head(10)
```

```
[29]:
```

	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	\
CustomerID							
18102	72	34	26632.62	4	4	4	

17949	70	32	22504.73	4	4	4
17450	70	28	18009.06	4	4	4
16029	80	39	15119.49	4	4	4
16013	70	24	10402.34	4	4	4
12901	81	20	5915.66	4	4	4
13798	72	34	4648.80	4	4	4
17857	72	12	4644.68	4	4	4
13694	71	32	4472.68	4	4	4
15061	73	23	3417.70	4	4	4

	RFMScore
CustomerID	
18102	444
17949	444
17450	444
16029	444
16013	444
12901	444
13798	444
17857	444
13694	444
15061	444

How many customers do we have in each segment?

```
[31]: print("Best Customers:␣
↪",len(rfm_segmentation[rfm_segmentation['RFMScore']=='444']))
print('Loyal Customers:␣
↪',len(rfm_segmentation[rfm_segmentation['F_Quartile']==4]))
print("Big Spenders: ",len(rfm_segmentation[rfm_segmentation['M_Quartile']==4]))
```

```
Best Customers: 218
Loyal Customers: 687
Big Spenders: 716
```